

Chapter 17 2nd Part

Making Complex Decisions

--- Decision-theoretic Agent Design


Xin Lu
11/04/2002



POMDP: UNCERTAINTY

- Uncertainty about the action outcome
- Uncertainty about the world state due to imperfect (partial) information

--- *Huang Hui*



Outline

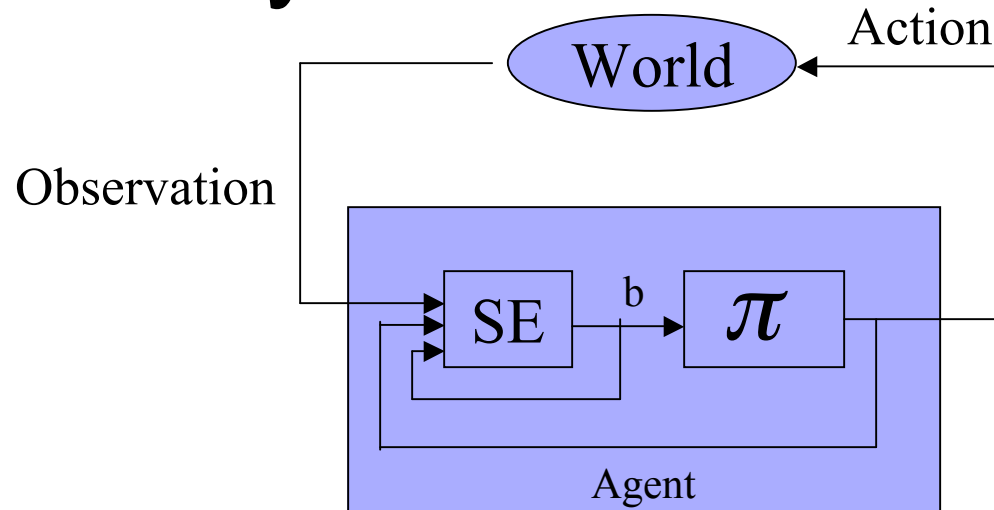
- POMDP agent




Constructing a new MDP in which the current probability distribution over states plays the role of the state variable causes the state new state space characterized by real-valued probabilities and *infinite*.

- Decision-theoretic Agent Design for POMDP

a limited lookahead using the technology of decision networks

Decision cycle of a POMDP agent



-  Given the current belief state b , execute the action $a = \pi^*(b)$
-  Receive observation o
-  Set the current belief state to $SE(b, a, o)$ and repeat.

Belief state

- $b(s)$ is the probability assigned to the actual state s by belief state b .

0.111	0.111	0.111	<u>0.000</u>
0.111		0.111	<u>0.000</u>
0.111	0.111	0.111	0.111

$$\left(\frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, 0, 0\right)$$

$$b'(s_j) = P(s_j | o, a, b) = \frac{P(o | s_j, a) \sum_{s_i \in S} P(s_j | s_i, a) b(s_i)}{\sum_{s_j \in S} P(o | s_j, a) \sum_{s_i \in S} P(s_j | s_i, a) b(s_i)} \longrightarrow b' = SE(b, a, o)$$

Belief MDP

- A belief MDP is a tuple $\langle B, A, \rho, P \rangle$:

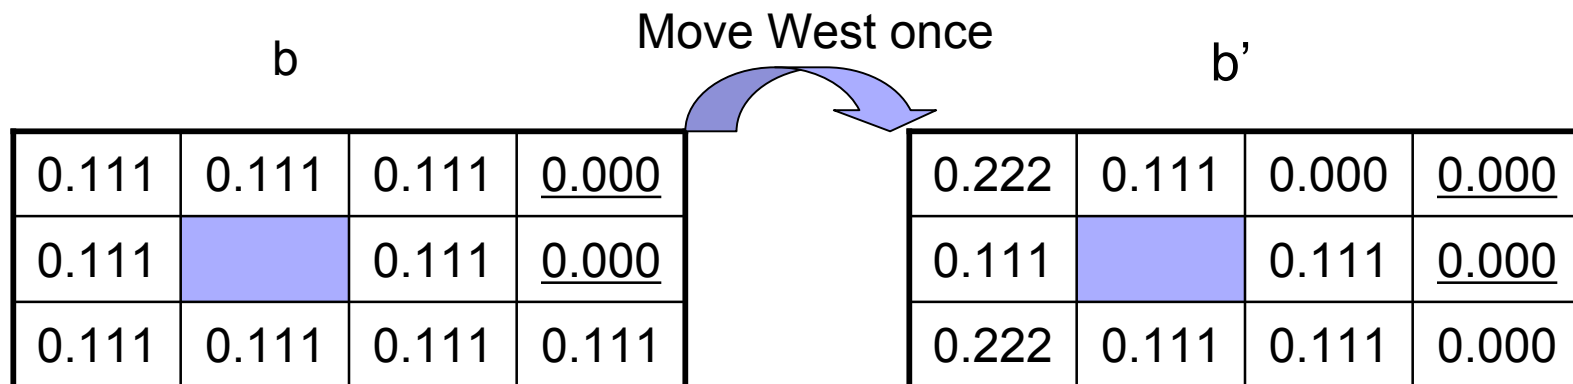
B = infinite set of belief states

A = finite set of actions

$$\rho(b, a) = \sum_{s \in S} b(s) R(s, a) \quad (\text{reward function})$$


$$P(b'|b, a) = \sum_{o \in O} P(b'|b, a, o) P(o|a, b) \quad (\text{transition function})$$

Where $P(b'|b, a, o) = 1$ if $SE(b, a, o) = b'$, $P(b'|b, a, o) = 0$ otherwise;



Solutions for POMDP

- Belief MDP has reduced POMDP to MDP, the MDP obtained has a continuous state space.
- Methods based on *value* and *policy iteration*:
A policy $\pi(b)$ can be represented as a set of *regions* of belief state space, each of which is associated with a particular optimal action. The value function associates a distinct *linear* function of b with each region. Each value or policy iteration step refines the boundaries of the regions and may introduce new regions.
- A Method based on lookahead search:
decision-theoretic agents






Decision Theory

= probability theory + utility theory

The fundamental idea of decision theory is that an agent is rational if and only if it chooses the action that yields the highest expected utility, averaged over all possible outcomes of the action.



A decision-theoretic agent

function DECISION-THEORETIC-AGENT(*percept*) returns *action*
 calculate updated probabilities for current state based on available
 evidence including current percept and previous action
 calculate outcome probabilities for actions
 given action descriptions and probabilities of current states
 select *action* with highest expected utility
 given probabilities of outcomes and utility information
 return *action*



Basic elements of decision-theoretic agent design

- Dynamic belief network--- the transition and observation models
- Dynamic decision network (DDN)--- decision and utility
- A filtering algorithm (e.g. Kalman filtering)---incorporate each new percept and action and update the belief state representation.
- Decisions are made by projecting forward possible action sequences and choosing the best action sequence.



Definition of Belief

- The belief about the state at time t is the probability distribution over the state given all available evidence:

$$Bel(X_t) = P(X_t | E_1...E_t, A_1...A_{t-1}) \quad (1)$$

X_t is *state variable*, refers the current state of the world

E_t is evidence variable.

Calculation of Belief (1)

- Assumption 1: the problem is *Markovian*,

$$P(X_t | X_1 \dots X_{t-1}, A_1 \dots A_{t-1}) = P(X_t | X_{t-1}, A_{t-1}) \quad (2)$$

- Assumption 2: each percept depends only on the state at the time

$$P(E_t | X_1 \dots X_t, A_1 \dots A_{t-1}, E_1 \dots E_{t-1}) = P(E_t | X_t) \quad (3)$$

- Assumption 3: the action taken depends only on the percepts the agent has received to date

$$P(A_{t-1} | A_1 \dots A_{t-2}, E_1 \dots E_{t-1}) = P(A_{t-1} | E_1 \dots E_{t-1}) \quad (4)$$



Calculation of Belief (2)

- Prediction phase:

$$\hat{Bel}(X_t) = \sum_{x_{t-1}} P(X_t | X_{t-1} = x_{t-1}, A_{t-1}) Bel(X_{t-1} = x_{t-1}) \quad (5)$$

x_{t-1} ranges over all possible values of the state variables X_{t-1}

- Estimation phase:

$$Bel(X_t) = \alpha P(E_t | X_t) \hat{Bel}(X_t) \quad (6)$$

α is a normalization constant



Design for a decision-theoretic Agent

Function DECISION-THEORETIC-AGENT(E_t) returns an action

inputs: E_t , the percept at time t

static: BN , a belief network with nodes X

$Bel(X)$, a vector of probabilities, updated over time

$$\hat{Bel}(X_t) \leftarrow \sum_{X_{t-1}} P(X_t | X_{t-1} = x_{t-1}, A_{t-1}) Bel(X_{t-1} = x_{t-1})$$

$$Bel(X_t) \leftarrow \alpha P(E_t | X_t) \hat{Bel}(X_t)$$

$$action \leftarrow \arg \max_{A_t} \sum_{X_t} [Bel(X_t = x_t) \sum_{X_{t+1}} P(X_{t+1} = x_{t+1} | X_t = x_t, A_t) U(x_{t+1})]$$


return *action*



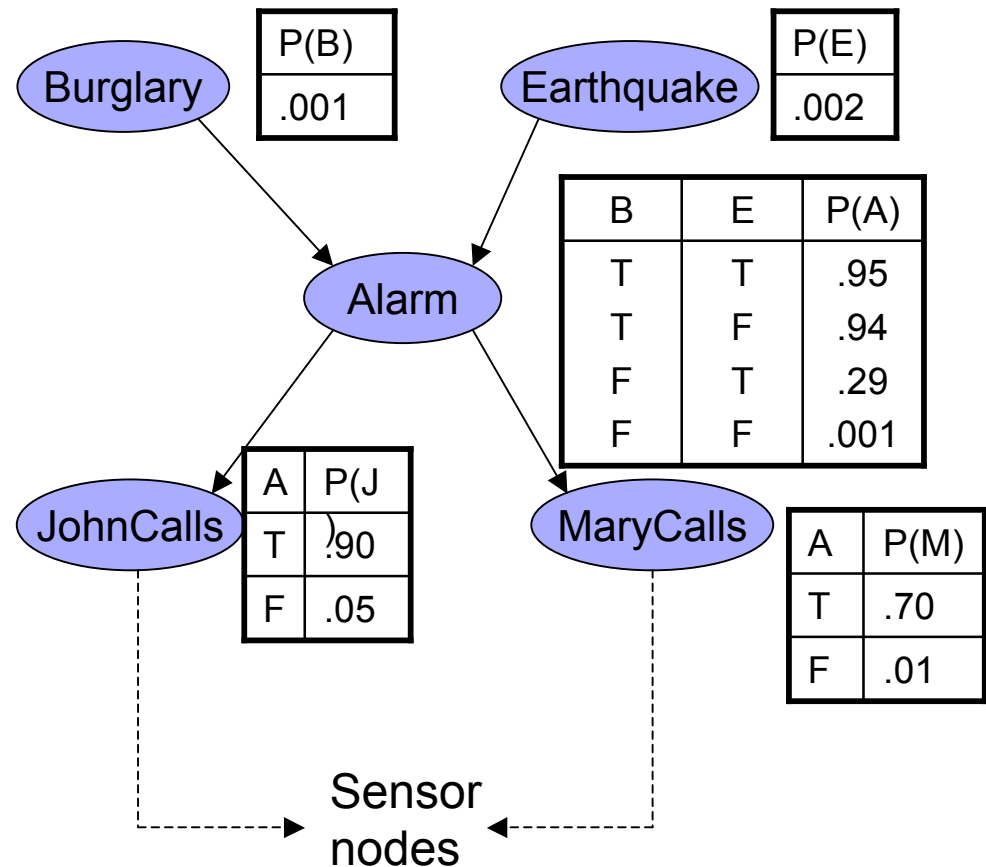
Sensing in uncertain worlds

- Sensor model: $P(E_t | X_t)$, describes how the environment generates the sensor data.
vs Observation model $O(s, o)$
- Action model: $P(X_t | X_{t-1}, A_{t-1})$, describes the effects of actions
vs Transition model $T(s, a, s')$
- Stationary sensor model: $\forall t \quad P(E_t | X_t) = P(E | X)$
where E and X are random variables ranging over percepts and states
Advantage: $P(E | X)$ can be used at each time step.


A sensor model in a belief network




(a) Belief network fragment showing the general relationship between state variables and sensor variables.



Next step: break apart the generalized state and sensor variables into their components.




(b) An example with pressure and temperature gauges



(c) Measuring temperature using two separate gauges

Sensor Failure

- In order for the system to handle sensor failure, the sensor model must include the possibility of failure.

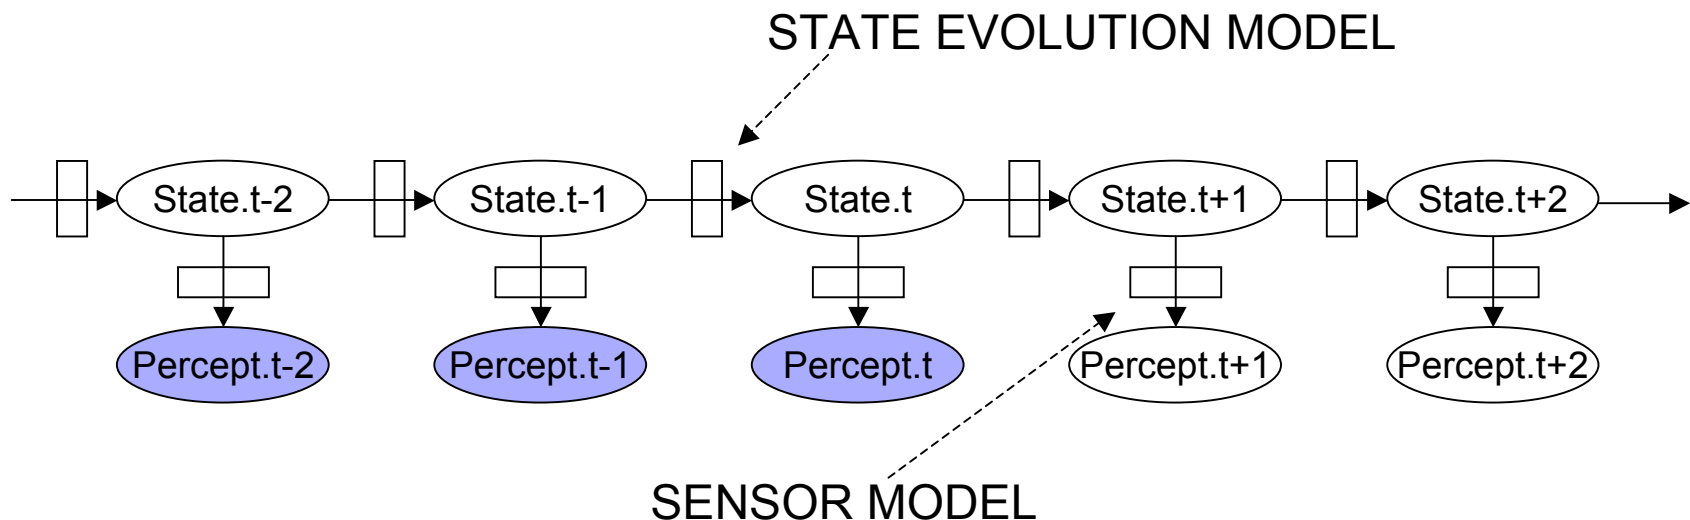




Dynamic Belief Network

- Markov chain (state evolution model):
a sequence of X_t values where each one is determined solely by the previous one: $P(X_t | X_{t-1})$
- Dynamic belief network (DBN): a belief network with one node for each state and sensor variable for each time step.

Generic structure of a dynamic belief network



Two tasks of the network:

- Calculate the probability distribution for state at time t
- Probabilistic projection: concern with how the state will evolve into the future

■ Prediction:

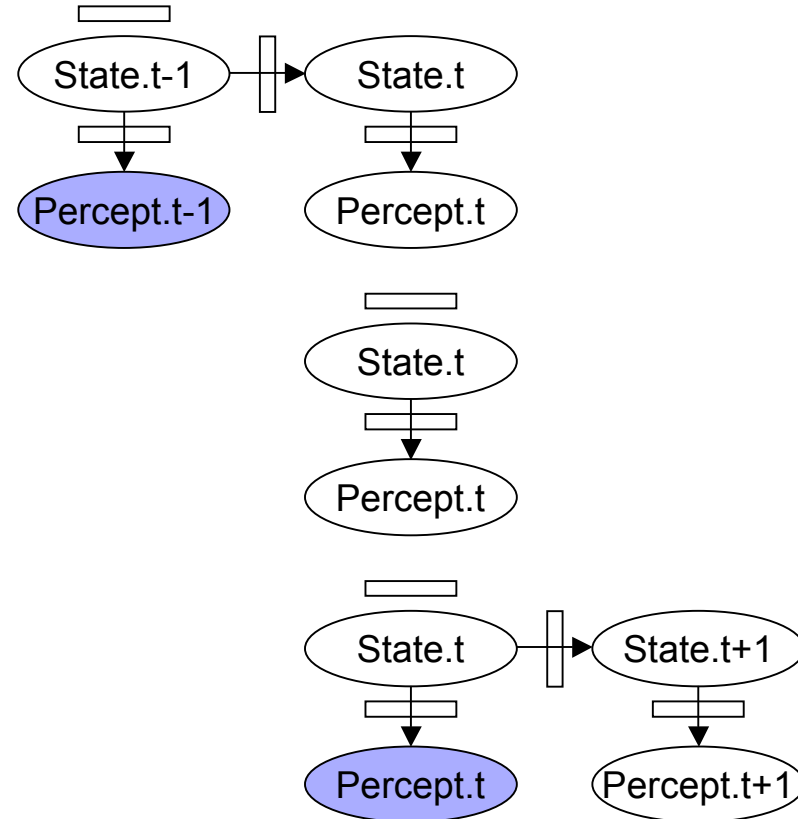
$$\hat{Bel}(X_t) = \sum_{X_{t-1}} P(X_t | X_{t-1} = x_{t-1}, A_{t-1}) Bel(X_{t-1} = x_{t-1})$$

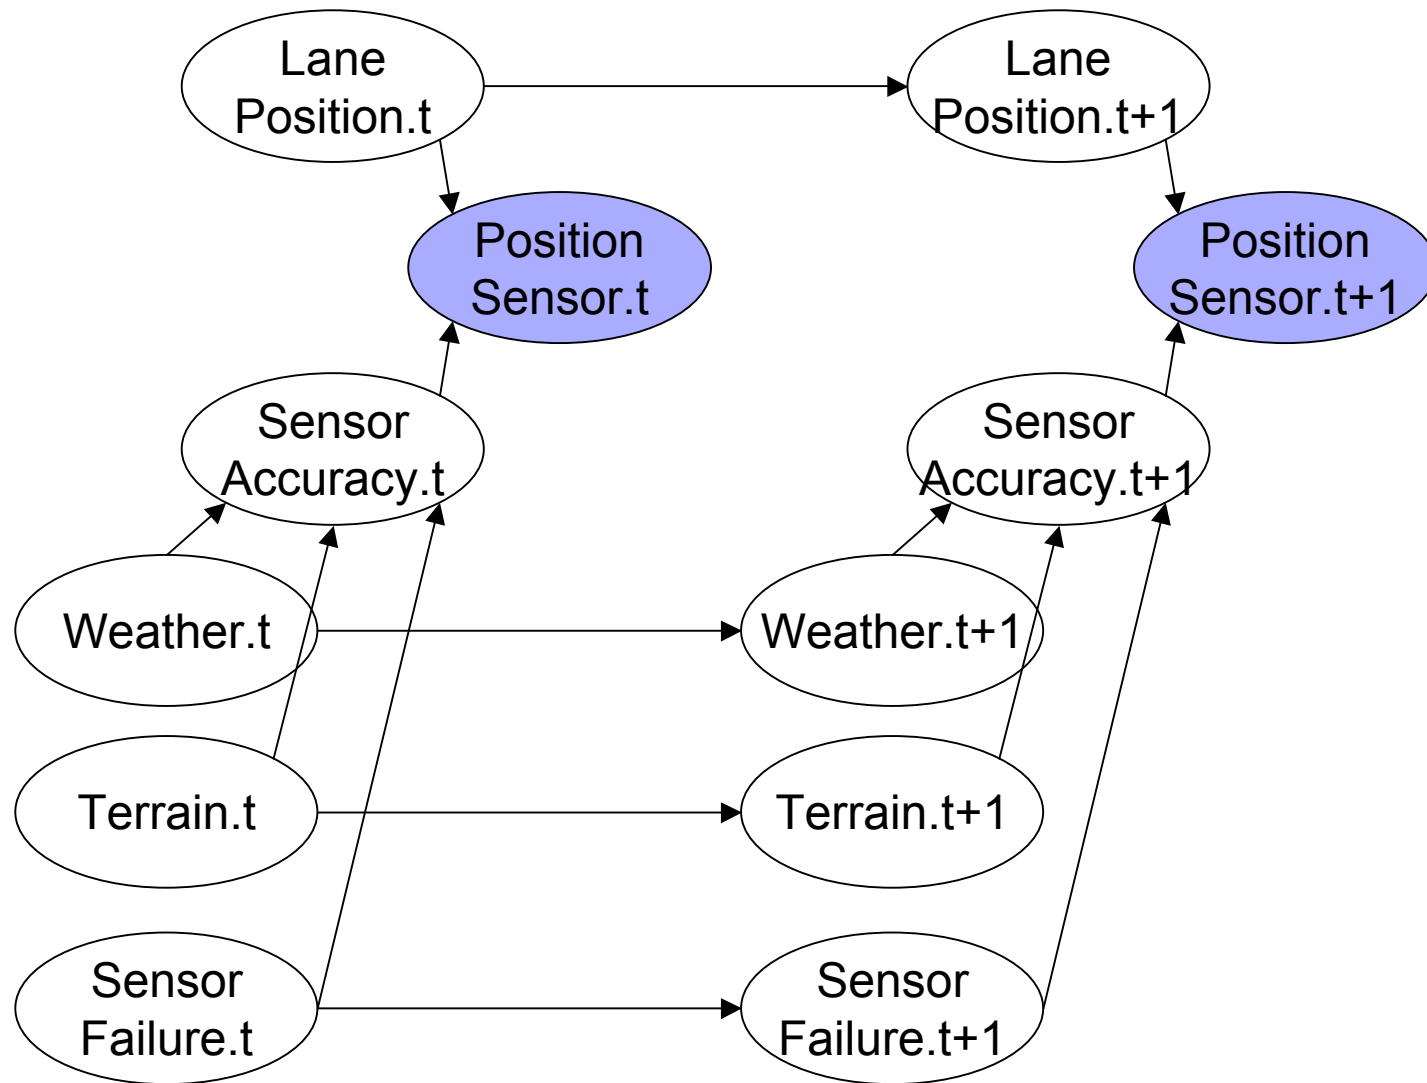
■ Rollup:

remove slice t-1

■ Estimation:

$$Bel(X_t) = \alpha P(E_t | X_t) \hat{Bel}(X_t)$$



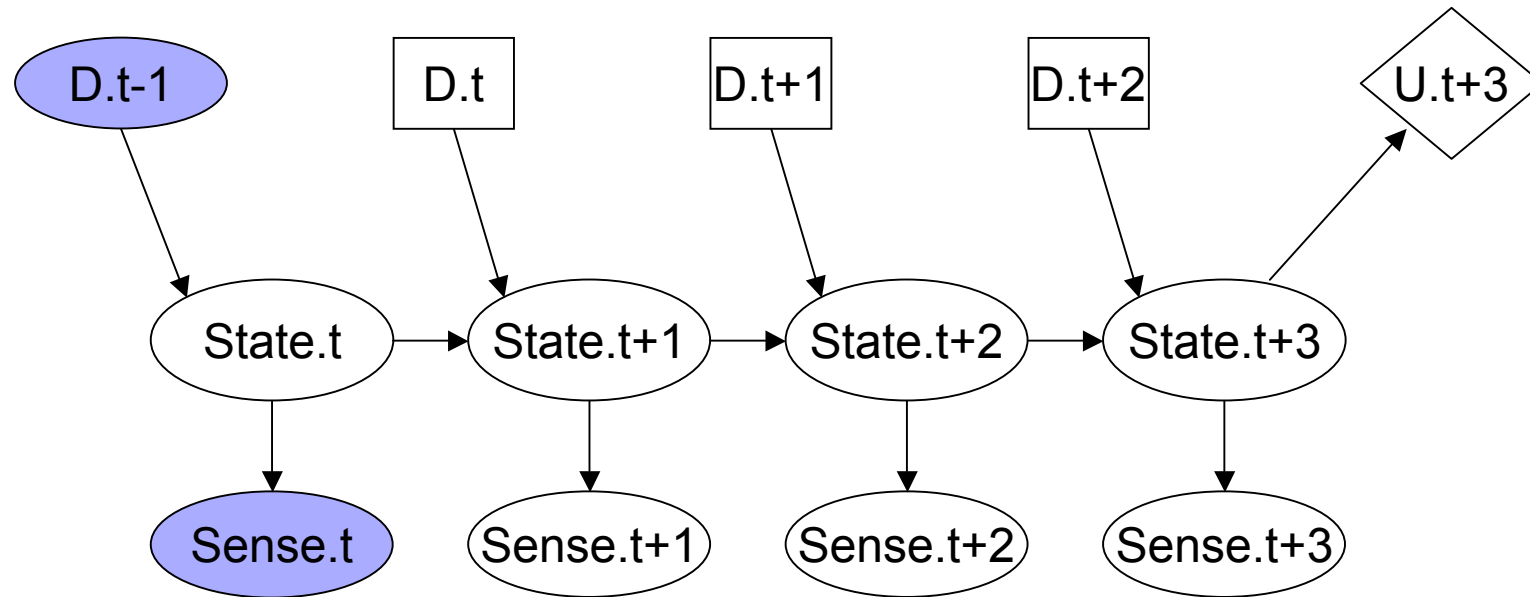




Dynamic Decision Networks


- Dynamic Decision Networks: add utility nodes and decision nodes for actions into dynamic belief networks.

The generic structure of a dynamic decision network






- The decision problem involves calculating the value of D_t that maximizes the agent's expected utility over the remaining state sequence.

Search tree of the lookahead DDN



Some characters of DDN search tree

- The search tree of DDN is very similar to the EXPECTIMINIMAX algorithm for game trees with chance nodes, expect that:
 -  There can also be rewards at non-leaf states
 -  The decision nodes correspond to belief states rather than actual states.
- The time complexity: $O(|D|^d \cdot |E|^d)$
 d is the depth, $|D|$ is the number of available actions, $|E|$ is the number of possible observations






Discussion of DDN

- The DDN promises potential solutions to many of the problems that arise as AI systems are moved from ***static, accessible***, and above all ***simple*** environments to ***dynamic, inaccessible, complex*** environments that are closer to the real world.
- The DDN provides a ***general, concise representation*** for large POMDP, so they can be used as inputs for any POMDP algorithm including value and policy iteration methods.



Perspective of DDN to reduce complexity

- Combined with a heuristic estimate for utility of the remaining steps
- Many approximation techniques:
 -  Using less detailed state variables for states in the distant future.
 -  Using a greedy heuristic search through the space of decision sequences.
 -  Assuming “most likely” values for future percept sequences rather than considering all possible values

...